

Aggregated Public Information Threat Model

Family privacy, address-doxxing risk, and non-phishing exposure assessment

Prepared for private family remediation use. Generated 5 June 2026. Redacted by design.

[Summary](#) [Address Risk](#) [No-Phishing Risks](#) [Misuse Matrix](#) [Child Privacy](#) [Developer Risk](#) [Actions](#)

[Queries](#) [Sources](#)

Private handling note: Do not upload this report publicly, place it in a public GitHub repository, or share it broadly. It intentionally avoids printing exact home addresses, exact child birth details, precise coordinates, private health details, raw financial values, child photos, or screenshots. It is a cleanup tool, not a new exposure document.

Table of contents

1. [Executive summary](#)
2. [Direct answer: can the exact current address be doxxed?](#)
3. [Direct answer: what risk remains if phishing links are never clicked?](#)
4. [Risk matrix of realistic misuse scenarios](#)
5. [Address-inference analysis](#)
6. [Child privacy analysis](#)
7. [Technical/developer exposure analysis](#)
8. [Social engineering and impersonation analysis](#)
9. [What to remove today](#)
10. [What to fix this week](#)
11. [What to monitor monthly/quarterly](#)
12. [What is probably okay to leave public](#)
13. [False positives and uncertainties](#)
14. [Search queries used](#)
15. [Removal and hardening checklist](#)

1. Executive summary

Bottom line: The main risk is not that one single page gives everything away. The risk is aggregation. Public dashboards, GitHub/Vercel apps, child-named apps, social snippets, old profiles, business/

professional records, and media metadata can be combined into a coherent map of the family, routines, interests, health/support context, work identity, and location history.

Overall risk rating: **HIGH** while the personal dashboards and child-named apps remain public. The risk drops substantially once those are removed, made private, renamed, and de-indexed.

Address doxxing

Exact current address confirmed in this pass: No. **Confidence:** Medium. The supplied audits did not confirm a direct current home address. However, exact address inference remains **possible** for a skilled person if they combine public records, old posts, image/video backgrounds, geotags, property/rental traces, and archived pages.

No phishing clicks

Not clicking links meaningfully reduces malware and credential-harvesting risk, but it does not stop phone impersonation, SMS/social DMs, account recovery guessing, SIM-swap attempts, identity verification abuse, reputational exposure, child profiling, or physical location inference.

Most urgent

Take down or protect the personal intelligence dashboard, JoshHub/life dashboard, child-named apps and repos, health/private-life apps, and any public child full-name/birth announcement snippets. Fix sources first, then search-engine caches.

Key findings

- **Public dashboards expose high-value context:** live checks confirmed public pages listing personal-data categories such as YouTube, health, ChatGPT, behaviour rhythms, finance, calendar, browser, locations, and history, and a family/life hub with finance, health, family-care, and support context. These are high-value aggregation surfaces.
- **Child-named apps create durable identity links:** public app/repo names and metadata tie children to parent developer accounts, age-stage context, routines, interests, and family identity. This is controllable and should be remediated first.
- **Exact address is not the only safety question:** suburb/town, workplaces, schools/community, routines, family roles, and service providers can still support impersonation or targeted harassment without a street address.
- **Old content matters:** old YouTube, Pinterest, Facebook snippets, academic profiles, old websites, GitHub commit history, files, PDFs, ZIPs, and web archives may preserve data after you change current profiles.

- **Security posture matters:** use phishing-resistant MFA where possible, unique random recovery answers, GitHub secret scanning, repository history review, Vercel Deployment Protection, and regular privacy reviews.

2. Direct answer: can the exact current address be doxxed?

Answer: From the supplied reports and current spot-checks, I would rate exact current address doxxing as **POSSIBLE BUT NOT CONFIRMED**. A casual stranger is unlikely to find the exact current home address from the visible sources alone. A skilled OSINT user has a better chance if they pursue old property/rental listings, public registries, image/video background clues, geotagged uploads, business records, cached pages, social comments, and archived versions of family apps or dashboards.

ACTOR LEVEL	COULD THEY GET EXACT ADDRESS?	COULD THEY NARROW LOCATION?	WHAT HELPS THEM?	WHAT IS MISSING / UNCERTAIN?	RATING
Casual stranger	Unlikely from confirmed sources alone	Yes - broad city/region and family/work context	Names, public profiles, dashboards, work/community hints	No confirmed public current street address in supplied audits	LOW-MEDIUM
Curious local person	Possible if they already know local context	Yes - workplaces, community, child/family routines may narrow practical location	Social snippets, work/school/community references, family apps, public media	Needs manual check of logged-out social posts, video backgrounds, old comments	MEDIUM
Skilled OSINT user	Possible	Yes - could build a strong location profile	Search operators, Git history, metadata, Wayback, ABN/business records, old property/rental traces, geotagged media	Needs complete search across archives and media, not just live pages	MEDIUM-HIGH
Persistent malicious actor using lawful public sources only	Possible but not guaranteed	Yes - especially if old address/context remains public	All of the above plus repeated monitoring and cross-platform correlation	Still constrained if sources are remediated quickly and no exact address appears in registries/posts	HIGH IMPACT / MEDIUM LIKELIHOOD

Important distinction: “Exact address not confirmed” is not the same as “safe.” The exposure is still meaningful because a person can build a believable family identity graph even without the street address.

3. Direct answer: what risk remains if phishing links are never clicked?

Never clicking phishing links blocks one common attack path. It does not remove the risk created by public information. A lot of misuse does not require you to click anything.

RISK REMAINS?	EXPLANATION	DOES NOT CLICKING LINKS STOP IT?
Phone/SMS impersonation	Someone can call or message while referencing real family, work, health, banking, app, or school context to sound legitimate.	Partly - you may not click links, but you might still trust the conversation.
Account recovery guessing	Public names, old schools, wedding/family context, hobbies, locations, and routines can help answer weak recovery prompts.	No.
SIM-swap or telco social engineering	Attackers may use biographical details to persuade a provider or to pass weak identity checks.	No.
Child profiling	Public child-first-name apps can reveal interests, routines, age-stage and parent identity.	No.
Location inference	Images, videos, geotags, registries, old posts, and dashboards can help infer place and routines.	No.
Reputation or embarrassment	Public health, finance, family-care, adult/private-life, or workplace apps may be misunderstood if found by employers, parents, students, or community members.	No.
GitHub/Vercel exposure	Repos and deployments can reveal metadata, app architecture, file paths, env variable names, and sometimes committed data.	No.
Old/cached content	Search engines and archives can keep snippets after you remove source content.	No.

4. Risk matrix of realistic misuse scenarios

SCENARIO	WHAT ENABLES IT	EFFORT	LIKELIHOOD	IMPACT	WARNING SIGNS	MITIGATION	DOES AVOIDING PHISHING LINKS PREVENT IT?
Exact address discovery	Old property/ rental listings, business records, social posts, geotags, image backgrounds, archive captures	Medium-High	Medium	Severe	People referencing street/home details; unfamiliar profile views; strange local contact	Remove geotags, audit media, clean public registries where possible, check archives, avoid house exterior posts	No
Suburb/town/ routine narrowing	Dashboards, school/work/ community references, family logistics, maps/ reviews, event posts	Low-Medium	High	High	Messages that know local routines or family schedule	Remove routine dashboards, avoid public pickup/ school/ church/ location patterns, delay posting	No
School/work/ health impersonation	Professional profiles, health/ support context, school/work references, names and roles	Low	High	High	Urgent messages about payroll, appointments, school, NDIS, devices, passwords, invoices	Verify using known phone numbers/ apps, never use contact details provided in the message, family code word	Partly

SCENARIO	WHAT ENABLES IT	EFFORT	LIKELIHOOD	IMPACT	WARNING SIGNS	MITIGATION	DOES AVOIDING PHISHING LINKS PREVENT IT?
Parent/relative impersonation	Family names, relationships, social graph, old posts, child routines/interests	Low-Medium	Medium	High	"Mum/Dad asked me to..." messages; requests for child pickup info or money	Family verification phrase; school pickup passwords; tell relatives not to post children publicly	No
Believable scam messages	Finance apps, banks mentioned, work history, GitHub/Vercel, AI/project context	Low	High	High	Messages referencing real banks, apps, repos, or recent projects	Out-of-band verification; password manager; MFA; strict payment/invoice approval rule	Partly
Security-question/account recovery guessing	Children's names, spouse name, school history, wedding/family context, hobbies, faith/community	Low	Medium	High	Account recovery emails/codes you did not request	Use random stored answers; remove public personal trivia; enable MFA/passkeys	No
Child identity correlation	Child-named apps, public full-name snippets, parent accounts, sibling/family references	Low	High while sources remain public	High	Child names appearing in search; strangers knowing interests or age-stage	Private/repos, neutral app names, no child full names, no public photos/routines	No

SCENARIO	WHAT ENABLES IT	EFFORT	LIKELIHOOD	IMPACT	WARNING SIGNS	MITIGATION	DOES AVOIDING PHISHING LINKS PREVENT IT?
Old photos/ videos found	YouTube, Pinterest, Instagram tags, Facebook snippets, community pages	Low-Medium	Medium	Medium-High	Old media surfaces in search/image results	Bulk-review videos/ photos; private family media; request tag/ caption edits	No
Cached/ archived page recovery	Search snippets, Wayback, GitHub caches, Vercel previews, old deployments	Medium	Medium	High	Old snippets remain after deletion	Fix source, submit Google/Bing refresh, request archive exclusions for controlled URLs	No
Behavioural profiling	Public personal intelligence dashboard, health/sleep/ finance/ YouTube/ ChatGPT categories	Low while public	High while public	High	Messages that understand habits, stressors, routines, interests	Take offline/ password-protect; delete raw datasets from public assets; de-index	No
Targeting public deployments	Public GitHub/ Vercel projects, environment variable names, source maps, raw assets	Medium	Medium	High if secrets/ data exist	Unexpected traffic, GitHub alerts, login attempts, Vercel logs spikes	Private repos, secret scanning, rotate keys, check git history and env files	No

SCENARIO	WHAT ENABLES IT	EFFORT	LIKELIHOOD	IMPACT	WARNING SIGNS	MITIGATION	DOES AVOIDING PHISHING LINKS PREVENT IT?
Reputational harm	Health, finance, family-care, relationship/intimacy, job-search, and personal dashboards	Low	Medium	High	Someone references a private app at work/church/school	Keep only professional portfolio public; use neutral naming and private deployments	No
Identity graphing across usernames	GitHub handles, old Wikipedia/YouTube/Pinterest/academic profiles, business name	Low	High	Medium	Search results cluster old identities together	Standardise public bios; delete unused profiles; separate public developer identity from family tools	No

5. Address-inference analysis

Confirmed in this pass

- No direct current home address was confirmed in the supplied child and spouse reports.
- The current web spot-check confirmed live public personal-data surfaces and child/family app surfaces, but this report does not use private conversation memory or print any address details.
- Some public sources can expose location-adjacent information: business/professional records, old profiles, geotagged media, event/community context, dashboard location-style categories, and social snippets.

Highest-value address-inference source categories

SOURCE CATEGORY	WHY IT MATTERS	ACTION
Public dashboards	Can reveal location-style categories, calendars, routines, finances, family-care context, and behaviour rhythms.	Take offline or password-protect immediately; confirm no raw location/CSV/JSON data is public.

SOURCE CATEGORY	WHY IT MATTERS	ACTION
Geotagged images/media	Precise metadata or recognisable backgrounds can reveal home/work/event locations.	Remove metadata, request deletion/suppression where justified, avoid uploading geotagged family media.
Video/photo backgrounds	House numbers, street views, car plates, uniforms, whiteboards, mail, appointment letters, or school signs can appear unintentionally.	Review old YouTube/Instagram/Facebook/Pinterest media while logged out.
Business/professional registers	May expose legal names, postcodes, business names, or official registration details.	Keep public business profile minimal; ensure official details are correct and not home-specific if alternatives exist.
Old property/rental/sale traces	Archived listings and social posts can connect a family to a house even after moving.	Search old and current address terms privately; request removals where possible.
GitHub files and history	Screenshots, local paths, PDFs, ZIPs, maps, CSV exports, and commits may contain addresses or coordinates.	Run local secret/data scans; make private repos private; rewrite history where necessary.
Search snippets and archives	Removed content can remain visible in snippets or archive captures.	Fix source first, then submit Google/Bing outdated-content refresh and Wayback exclusion requests for controlled URLs.

6. Child privacy analysis

The child-specific exposure is one of the clearest controllable risks. The issue is not that the apps are bad. The issue is that public URLs, app names, metadata, READMEs, and social snippets connect children to parent identities, family structure, age-stage details, routines, and interests.

EXPOSURE TYPE	REVEALS	URGENCY	ACTION
Public child full-name or birth-related social snippets	Child identity, parent identity, family relationship, exact birth-context clues	REMOVE TODAY	Delete/restrict original posts and duplicates; request search cache refresh.
Child-named GitHub repos and Vercel/GitHub Pages apps	First name, parent developer account, interests, age-stage, routines, learning/play context	REMOVE TODAY	Make private or delete; rename to neutral project names; disable public deployments.
Bedtime/media/routine apps or files	Family routine, sibling context, media habits, timing clues	REMOVE TODAY	Private repo/deployment; remove watch-history exports and ZIPs from public history.
Parent social bios listing children	Family structure and child names tied to parent identities	THIS WEEK	Use initials or neutral phrase; enable tag review.

EXPOSURE TYPE	REVEALS	URGENCY	ACTION
Videos/photos of children or home	Faces, voices, school uniforms, location clues, house/car details	THIS WEEK	Make family videos private; review thumbnails and captions; ask third parties to edit tags.
Health/therapy/NDIS/development notes	Sensitive child support context	NEVER PUBLIC	Keep out of public repos, dashboards, app prompts, screenshots, and social captions.
Generic demo app using fictional data	Educational idea without family identifiers	ACCEPTABLE IF ANONYMISED	Use fictional names, no photos, no locations, no personal routines, no indexable child metadata.

Family rule to adopt

Public child data rule: no child full names, exact birthdays, school/preschool names, health/developmental details, therapy/NDIS details, public face photos, home exterior images, or routine/location clues unless both parents deliberately agree and understand the risk.

7. Technical/developer exposure analysis

The developer footprint is valuable because it can expose both personal identity and technical attack surface. Public source code does not need to contain a password to be useful to an attacker; file paths, environment variable names, app architecture, deployment names, screenshots, test data, PDFs, and raw exports can still help.

AREA	RISK	RECOMMENDED CONTROL
Personal dashboards	Potential public raw JSON/CSV/source maps; life-map and routine profiling	Take offline; make repo private; inspect build output; remove raw assets; de-index.
GitHub repos	Child names, health, finance, work, private-life apps, commit history, local paths	Sort into public portfolio vs private life; make all personal/family/health/private-life repos private.
Secrets and credentials	Even no confirmed live key is not proof; history and branches need scanning	Enable GitHub secret scanning, run local gitleaks/trufflehog, rotate anything exposed.
Vercel deployments	Old previews, production URLs, app metadata, public build assets	Use Deployment Protection/password protection where appropriate; delete old deployments; review env vars.
File uploads/assets	PDFs, ZIPs, screenshots, media, reports may contain sensitive data	Search repos for PDFs, ZIPs, screenshots, CSV/JSON, addresses, names, medical/finance terms.

AREA	RISK	RECOMMENDED CONTROL
Public portfolio	Legitimate professional benefit, but must be separated from family life	Keep a clean public GitHub with only sanitised portfolio projects and generic demo data.

8. Social engineering and impersonation analysis

Public family context increases the believability of scams. This matters even if no one clicks links. Voice calls, texts, DMs, fake invoices, fake school/health messages, and “I know your family” approaches can all be built from public context.

Likely impersonation themes

- **School/work IT:** fake Microsoft/Google password, device, payroll, onboarding, or ticket requests.
- **Healthcare/NDIS/family support:** fake appointment, invoice, provider, health-record, or plan-management messages.
- **Banking/budgeting:** fake Bendigo/ANZ/YNAB-style messages or calls.
- **GitHub/Vercel/developer services:** fake security alerts, deployment errors, repo warnings, or domain-renewal notices.
- **Family/community:** messages referencing children, spouse, church/community, old towns, old employers, or relatives.

Family verification practice

Use an out-of-band rule: If any message creates urgency around money, passwords, children, health, school pickup, account access, or identity documents, stop and verify through a known-good channel. Do not use the phone number, link, or email supplied in the message itself. Use a family code word for child/family requests.

9. What to remove today

1. **Take the personal intelligence dashboard offline or put it behind strong authentication.** Check for raw CSV/JSON/build assets and source maps.
2. **Make the life/family dashboard private or delete the production deployment.** Remove finance, health, family-care, spouse, and child support details from public surfaces.
3. **Make child-named repos private immediately.** Disable GitHub Pages/Vercel deployments or protect them.
4. **Rename child-related app metadata.** Change title, manifest name, package name, icon filenames, README headings, repo names, route names, and social preview metadata.
5. **Remove/restrict child full-name or birth announcement posts and duplicates.** Then request search-result refresh.

6. **Private/delete health, finance, adult/private-life, job-search, and family-care apps unless fully anonymised.**
 7. **Enable GitHub secret scanning and run a local history scan.** Rotate any exposed credentials even if you later delete the repo.
 8. **Review public YouTube/Pinterest/Facebook/Instagram content for child faces, home details, car plates, school uniforms, documents, and location clues.**
-

10. What to fix this week

1. Create a clean public/private split: one public professional portfolio; all family/health/work-internal/private-life apps private.
 2. Search logged out in a private window for parent names, handles, child first-name + parent-name combinations, old towns, business name, and app names.
 3. Request edits/removal from third-party posts that expose child names, parent tags, family relationships, health-support context, or location clues.
 4. Use Google Refresh Outdated Content after source changes. Use Bing removal/refresh options where needed.
 5. Check Wayback Machine captures for controlled GitHub Pages/Vercel/custom-domain URLs and request exclusions where appropriate.
 6. Replace weak security-question answers with random strings stored in a password manager.
 7. Enable MFA/passkeys on Google, Microsoft, GitHub, Vercel, bank, telco, Facebook/Instagram, YouTube, and domain registrar accounts.
 8. Update public LinkedIn/professional bios to a minimal current professional version without family/private detail.
-

11. What to monitor monthly/quarterly

- Monthly: Google/Bing searches for name variants, handles, business name, and old app names.
- Monthly: GitHub/Vercel projects list - confirm no personal app is public by accident.
- Quarterly: child privacy audit across social platforms, GitHub, Vercel, YouTube, Google Drive sharing links, old websites, and web archives.
- Quarterly: password manager and MFA audit, especially account recovery methods and telco protections.
- After every new app: check title, manifest, README, package name, deployment URL, metadata, screenshots, PDF/ZIP assets, source maps, and noindex settings.

12. What is probably okay to leave public

- A minimal professional profile with your preferred name, broad skills, and business/professional contact route.
- Sanitised portfolio apps using fictional data and neutral project names.
- Generic D&D/games/learning demos that do not include family names, private routines, health, finances, addresses, workplaces, schools, or raw exports.
- Official register entries that cannot reasonably be removed, as long as they do not expose unnecessary private contact details.

13. False positives and uncertainties

- **Not found is not proof of absence.** The supplied reports did not confirm a current home address, private phone number, live credentials, or private email in the sampled public audit, but public data changes constantly and some platforms only expose content to logged-in users or snippets.
- **Some social snippets require manual verification.** Instagram/Facebook/YouTube/Pinterest may show different content when logged in, logged out, on mobile, or via search snippets.
- **Name collisions exist.** Do not contact unrelated people or attempt to remove content that is not yours.
- **Private conversation memory was not used as a public source.** This report evaluates public-source risk and supplied audits, not private knowledge from chats.
- **Repository history was not fully cloned and scanned in this report.** Treat secret/data exposure status as unconfirmed until local tools scan all repos, branches, releases, assets, and commit history.

14. Search queries used

To avoid turning this report into a search blueprint with sensitive exact strings, the queries below are redacted patterns rather than exact child/full-name searches.

PURPOSE	REDACTED QUERY PATTERN
Family identity graph	"[Parent full name]", "[Parent preferred name]" + "[spouse first name]", "[business name]"
Child exposure	"[child first name]" + "[parent handle]", "[child app name]", site:github.com [parent handle] [child app]
Developer exposure	site:github.com [handle] .env, site:github.com [handle] "C:\\Users", site:github.com [handle] "OneDrive", site:github.com [handle] filetype:pdf
Deployment exposure	site:vercel.app [handle or app name], site:github.io [handle or app name]

PURPOSE	REDACTED QUERY PATTERN
Media/location risk	"[parent name]" "YouTube", "[parent name]" "Pinterest", "[parent name]" "Wikimedia Commons"
Cached/archive checks	cache/search snippet checks for removed URLs , Wayback Machine checks for controlled URLs
Official remediation guidance	GitHub remove sensitive data repository , Vercel Deployment Protection , Google Refresh Outdated Content , Wayback Machine removal request

15. Removal and hardening checklist

PLATFORM / AREA	CHECKLIST
GitHub	Make private repos private. Disable Pages where not needed. Review branches/releases/issues/assets. Enable secret scanning. Remove sensitive history carefully if needed. Rotate any exposed secrets.
Vercel	Delete/protect family/private deployments. Enable Deployment Protection/password protection where appropriate. Delete old previews. Review env vars and build outputs.
Google Search	Fix the source first. Then use Refresh Outdated Content or Search Console removal options as appropriate.
Bing	Use Bing Webmaster/removal tools after the source page is changed or removed.
Internet Archive	Check captures of URLs you controlled. Request exclusion/removal for controlled site/account URLs where justified.
Instagram/Facebook	Remove child names from bios. Enable tag review. Archive/delete public child/family posts. Ask third-party posters to edit captions/tags.
YouTube	Set family/home/child videos to Private, not merely Unlisted. Review thumbnails, descriptions, playlists, and old Shorts.
Pinterest	Make personal/family/wedding/health/location boards private. Remove full-name couple/family pins if not needed.
Accounts	Use password manager, unique passphrases, MFA/passkeys, random recovery answers, and secure telco account PINs.
Family process	Adopt a child privacy rule, a family verification phrase, and an out-of-band verification rule for money/identity/children/health requests.

Sources and evidence basis

Uploaded source reports reviewed

- Defensive Public Exposure Audit - Joshua / family public exposure report, 5 June 2026.
- Child Privacy OSINT Audit - Child 1, redacted report, 5 June 2026.
- Child Privacy OSINT Audit - Child 2, redacted report, 5 June 2026.
- Defensive OSINT Public Exposure Audit - Kristy, redacted report, 5 June 2026.

Current live spot-checks performed

- Public personal intelligence dashboard - confirmed accessible and listing personal-data categories, 5 June 2026.
- Public family/life hub - confirmed accessible and listing finance, health, family-care, and support context, 5 June 2026.
- Public child phonics app - confirmed accessible and containing child-name, literacy-age, parent-coach, microphone/speech and region context, 5 June 2026.
- Public child app GitHub repositories - confirmed accessible and containing README/app metadata connecting child-first-name branding with parent developer accounts, 5 June 2026.

Official remediation and safety references

- eSafety Commissioner - Doxing guidance and reporting online harm: [eSafety doxing guidance](#) and [reporting page](#).
- OAIC - personal information, data breaches, and identity fraud guidance: [what is personal information](#), [identity fraud](#).
- Scamwatch - phishing, impersonation and scam reporting guidance: [phishing scams](#), [report a scam](#).
- Australian Signals Directorate / cyber.gov.au - MFA and small business cyber security guidance: [small business cyber security](#).
- GitHub Docs - secret scanning and removing sensitive data from a repository: [secret scanning](#), [remove sensitive data](#).
- Vercel Docs - Deployment Protection and password protection: [Deployment Protection](#).
- Google Search Help - Refresh Outdated Content / remove web results from Google Search: [Refresh Outdated Content](#), [remove web results](#).
- Internet Archive Help - request exclusion/removal from archive.org: [removal request guidance](#).