

QA & UX Audit

DCSPrep Professional Development Apps

Subjects	dcspd.vercel.app dcs-professional-development.vercel.app
Audit framework	Five-pillar IT-education QA — adapted for personal-use scope
Reviewer role	Senior QA Engineer & UX Auditor (IT education and PD)
Date prepared	Friday, 8 May 2026
Audience	Josh Parris — product owner and primary user

Executive summary

Both apps share the DCSPrep brand, identical privacy and operational-priority guardrails, and a strong question-first pedagogy that's well-fitted to a Level 1 school-IT context. However, the two apps are running in parallel under the same name with different feature scopes and different working states — **dcsp-professional-development** is the more focused, fully functional product, while **dcspd** has a much broader feature surface but appears to have a critical issue rendering individual module pages, which alone makes it unsafe to use as the primary tool until verified.

Headline recommendation

Pick one app as canonical for the next quarter. **dcsp-professional-development** is the safer choice today because its module flow works end-to-end. Park **dcspd** as a feature-experiment branch and only promote features once their pages render reliably and the navigation is consolidated.

A note on the audit framework

The brief asked for a five-pillar audit designed for commercial IT-education platforms — covering Course Schema for SEO, instructor credentials, partner badges (AWS, CompTIA), student testimonials, enrollment funnels, and conversion. Several of those pillars don't map onto what these apps actually are.

Both DCSPrep apps are **personal professional-development tools** for one user, in a school IT context, with explicit on-page warnings that they're for personal PD only and must not contain real student or staff data. There are no enrolments, no instructors, no partner badges to display, and no SEO conversion goals — so applying those criteria literally would produce noise rather than signal.

I've kept the five-pillar structure but reframed each pillar against the real job these apps do: helping a single L1 tech build durable, evidence-backed support skills during quiet windows. Where a pillar genuinely doesn't apply, I say so plainly rather than inventing findings to fill the section.

App comparison at a glance

Dimension	dcspd.vercel.app ("App A")	dcs-professional-development.vercel.app ("App B")
Brand & title	DCSPrep	DCSPrep (identical)
Module count	32 modules	36 modules (incl. RBC + SMITB)
Module detail page	Empty body in server response	Working — tabbed Start / Learn / Review / Assessment
Top-nav surface	~16 links (dashboard); ~27 links on inner pages	~13 links, consistent across pages
Unique features	Skill Coach, Network Map, Classroom Desk, Roleplay Bot, Voice-to-Ticket, Hardware, Certificates, KB Sync, Supervisor, Mobile QA, Admin Hub, Sync, Trainer Guide	Progress page, Support Tools, light/dark toggle, "Overwhelmed mode" tile
Pedagogy	Daily Mastery prompt, scenario lab, AI mentor 7-day plan	Diagnostic-first per module, "tiny" due-today mode, scenario lab
Operational guardrails	Prominent on every page	Prominent on every page
Privacy banner	Header + footer	Footer only
Theme toggle	Not visible	"Light" toggle in header
Empty-state stats	"0 days · Best: 0 days"	0% / 0 / 0 min — less judgemental wording

Reading this table

App A is the more ambitious vision — more modules-adjacent features, AI mentor flows, simulations. App B is the more disciplined product — tighter scope, working module pages, cleaner navigation. The right move depends on whether you're ready to ship features or still iterating on the learning core.

Critical fixes (Priority 1)

Immediate blockers — resolve before further feature work.

P1

App A: individual module pages return effectively empty server content

Finding	When fetching /modules/ports-and-protocols and /modules/dcs-it-support-foundations on App A, the server response contains the navigation chrome, the “Support comes first” banner, and the privacy footer — and nothing else. There is no module title, no Start / Learn / Review / Assessment flow, no diagnostic prompts. The same routes on App B return fully formed module pages.
Evidence	Likely cause: client-side-only rendering of module bodies, with no SSR or static fallback — or a real runtime error that’s swallowed and shows a blank pane. Either way, a learner who lands on a module is left with nothing to do, and any non-JS surface (search engines, link previews, archived snapshots, screen readers in some configurations) sees an empty page.
Recommended fix	Verify in a real browser with DevTools open: is content actually rendering client-side and just absent from server HTML, or is the page genuinely blank? If client-only, add server rendering or a static fallback for module bodies. If genuinely broken, check the build — the module list page works, so the data exists; the route handler or page component is the suspect.

P1

Two production deployments share the same brand and confuse the source of truth

Finding	Both apps title themselves “DCSPrep”, share the same operational guardrails, the same module set (mostly), and the same footer copy. From the outside it’s impossible to tell which is canonical, which is a staging branch, and which is being actively maintained.
Evidence	If you open one on a phone and the other on a laptop, your progress, PD log, and feedback log are split across two stores you may not even realise are separate. For a tool whose value depends on accumulated evidence and streaks, this fragments the very data the dashboard depends on.
Recommended fix	Pick one as canonical. Redirect the other’s root to it (or take it offline) and put a single line on the surviving app: “Replaces dcsdpd.vercel.app — v2” or similar, with the date. If both must coexist (e.g. one is a sandbox), rename the brand on the secondary to “DCSPrep Sandbox” so the distinction is visible at a glance.

P1 Empty-state numbers undermine first-run confidence	
Finding	On the App A dashboard, the headline stats read: <i>Learning streak 0 days · Best: 0 days · Overall progress 0% · Due today 0 · This month logged PD 0 min · 0 scenarios</i> . App B is a touch gentler but still leads with mostly zeros.
Evidence	For a tool whose explicit aim is steady, sustainable growth, the front door currently reads “you have done nothing.” That’s a poor framing for a user with ADHD trying to keep momentum, and it makes the app feel broken even when it isn’t.
Recommended fix	Replace zero-state metrics with task-shaped first actions: “Start your first 5-minute review”, “Set a streak goal”, “Pick today’s focus topic.” Hide “Best: 0 days” entirely until a streak has ever been set. Show stats only once they have content.

UX & learning-flow optimisations (Priority 2)

Improvements to the learner journey — not blockers, but high-value.

P2

App A's navigation surface is too wide for a single-user tool

Finding

Inner pages on App A expose 27 top-level navigation links: Dashboard, Skill Coach, Focus, Modules, A+ Core 2, Scenario Lab, Network Map, Classroom Desk, Roleplay Bot, Voice-to-Ticket, Hardware, Search, Feedback Log, Due Today, Strict Quiz, PD Log, Certificates, Academic PD, Evidence Pack, KB Lab, KB Sync, Error Log, Readiness, Supervisor, Mobile QA, Trainer Guide, Admin Hub, Sync, Settings.

Evidence

Twenty-seven sibling links is past the point where the eye can scan; on a phone it's a wall. Even for a power user, the cognitive load of choosing a next action goes up sharply when the choice space is this flat.

Recommended fix

Group into 5–6 sections with progressive disclosure: **Learn** (Modules, Academic PD, Skill Coach), **Practice** (Scenario Lab, Strict Quiz, Roleplay Bot, Network Map, Classroom Desk, Voice-to-Ticket), **Track** (Focus, Due Today, PD Log, Evidence Pack, Readiness, Certificates, Feedback Log), **Reference** (Search, KB Lab, Hardware, Trainer Guide, Error Log), **Admin** (Settings, Sync, KB Sync, Admin Hub, Mobile QA, Supervisor). Hide Admin behind a settings panel rather than the main nav.

P2

Both apps render duplicate navigation in the page body

Finding

On the App A dashboard, the full navigation list appears **twice** in close succession — once as the header and once again as a secondary list. App B does the same with its main nav and a labelled “Navigation” sidebar listing the same destinations.

Evidence

A user reading top-down sees the same list of links twice in a few seconds. For screen-reader users this means navigating past the menu twice. It also adds noise that makes the actual content of the page harder to find on small screens.

Recommended fix

Render the navigation once. If a secondary nav is genuinely needed (e.g. a contextual side-rail on inner pages), make sure it lists *different* links — section-specific shortcuts, not the same global menu.

P2

“Last updated” signals are missing across all technical content**Finding**

Module cards show category, level, and estimated minutes — but no last-edited or last-verified date. A learner cannot tell whether the Intune content was reviewed last month or last year.

Evidence

School IT moves: Entra renames, Intune capability shifts, ViewBoard firmware behaviour changes. Without a date, every claim has to be re-verified before use, which slows the very thing the app exists to speed up.

Recommended fix

Add a *Last reviewed: 8 May 2026* line to every module page (App B already has the structure for this in the module header). Surface the same date as a small badge on module cards. When a module hasn't been touched in 6+ months, colour the badge amber as a soft prompt to re-verify.

P2

Daily Mastery prompt overlaps with module content**Finding**

App A's dashboard shows the Daily Mastery prompt *“Explain why remembering a few common ports is useful to a DCS Level 1 tech even if Josh should not change firewall rules.”* The exact same prompt appears as Prompt 2 inside the App B Ports and Protocols module.

Evidence

If Daily Mastery is meant to be a daily retrieval cue, drawing from the same pool the modules use is fine — but a learner shouldn't see the same prompt twice in a session and assume one is the answer key. There's also a third-person reference to “Josh” baked into a question that the app shows to Josh, which reads slightly oddly.

Recommended fix

Either tag prompts so a module's own prompts are excluded from that day's Daily Mastery while the user is in that module's window, or rephrase Daily Mastery items in second person (*“Explain why remembering common ports is useful to you as a Level 1 tech...”*). Pick one voice and apply it consistently.

P2

App A's “Quick actions” bar is a list of equally weighted links**Finding**

Eleven equally weighted action buttons sit under “Quick actions”: Start tiny, 10-question assessment, Scenario Lab, Review flashcards, Add PD log, View readiness, Learning Paths, Feedback Log, Global Search, Technical Playground, Open evidence pack.

Evidence

“Quick” implies a single high-priority next action. Eleven equal options is a menu, not a quick action.

Recommended fix

Show the top one or two actions large (driven by the user's actual state — due reviews, weak topic, last action), and demote the rest to a collapsible “More actions” group. App B's “Three tiny moves when everything feels loud” pattern is closer to the right shape — consider porting that mental model.

P2 “Overall progress 0%” mixes incommensurable things	
Finding	The progress metric is described as “Across modules, flashcards, practical outputs, and assessment sessions” — four very different units smashed into a single percentage.
Evidence	A user can’t tell whether 30% means “you’ve finished 30% of modules” or “you’re due 30% on flashcards.” Composite metrics that hide their components tend to either flatline or whiplash, and either erodes trust in the number.
Recommended fix	Either show four small bars (one per dimension) or pick a single primary metric for the dashboard headline and move the others to a Progress detail page. App B already has a /progress route — that’s the natural home for the breakdown.

Polishing (Priority 3)

Minor copy, style, and consistency fixes — nice-to-haves.

P3 Tonal inconsistency in module subtitles

Finding	Several module subtitles slip into mildly defensive phrasing: “... <i>without pretending Level 1 should edit firewall policy</i> ”, “... <i>without overpromising access</i> ”, “... <i>without ever collecting passwords</i> .” The intent is good (scope clarity) but the phrasing is slightly anxious in tone.
Evidence	These read as guarding against a misuse the learner wasn’t about to commit. Replace with positive scope language: “ <i>Diagnose port-related symptoms and describe them clearly to L2</i> ” or “ <i>First-line password support, with self-service reset and lockout patterns</i> .” Same scope, calmer voice.
Recommended fix	Sweep all module subtitles in one pass. If a phrase starts with “without pretending”, “without ever”, or “without overpromising”, rewrite as a positive statement of what the learner <i>does</i> learn.

P3 Footer privacy reminder duplicates a banner already at the top

Finding	On most pages, the privacy reminder appears twice — once near the header (“This app is for personal PD only. Do not enter real student, staff, parent, credential, or network details.”) and again in the footer (“DCSPrep is a personal PD tool. Keep all entries privacy-safe...”).
Evidence	The second instance has diminishing returns. By the time a user has scrolled past it five times in a session, they stop reading it — including the times when it actually matters (e.g. when typing into a free-text field).
Recommended fix	Keep the footer line as a quiet reminder, but trigger a contextual nudge at the moment a user opens any free-text input that could capture sensitive data — PD log entry, error log, ticket-note simulator. That’s where the warning has bite.

P3 Mixed punctuation in tag chips

Finding	Module tag chips on the modules list page are run together with no separator: “ <i>portsprotocolsfirewall thinkingnetwork basics</i> ”, “ <i>DNSDHCPgatewayIP169.254SSIDBYODforget rejoincross-device comparison</i> .”
Evidence	Hard to scan, hard to scan-read for screen readers, and looks like a styling regression rather than a deliberate choice.
Recommended fix	Render tags as proper chips with explicit gaps and aria-labels (“ports”, “protocols”, “firewall thinking”), or fall back to comma separation (“ports, protocols, firewall thinking, network basics”). Either reads cleanly.

P3

App B: “20-minute focus block” quick action links to /modules**Finding**

On the App B dashboard, the action labelled “20-minute focus block — Use the main guided study block for one bounded session” targets **/modules**. App A has a dedicated **/focus** route for this purpose.

Evidence

The label promises a bounded 20-minute session; the destination is an open module list. Mismatch between promise and destination erodes trust in the rest of the dashboard’s actions.

Recommended fix

Either point this action at a real focus-mode route (port the App A pattern), or rename the action to “Pick a 20-minute module” so the destination matches the label.

P3

Heading length varies across module cards**Finding**

Module titles range from 3 words (“Ports and Protocols”) to 10+ words (“A+ Core 1: Troubleshooting (hardware + network + printers)”). Some carry a prefix (“A+ Core 1:” or “RBC” or “SMITB”), some don’t.

Evidence

On a card grid, varied lengths make scanning slower. Long parenthetical scopes also wrap in an ugly way on narrow screens.

Recommended fix

Standardise: short title (≤6 words), with course-prefix moved to the category badge area, and parenthetical scope demoted to the description. So “*Troubleshooting: hardware, network, printers*” with an “A+ Core 1” badge.

Pillar-by-pillar audit

The brief specified five pillars. Below, each is applied honestly to a personal PD tool — with sub-criteria marked as *applicable*, *partial*, or *not applicable* where the framework was built for a different kind of product.

1. Technical reliability & SEO

Sub-criterion	Status	Note
Broken links	OK	All sampled internal links resolved; no 404s observed in scope.
Slow-loading descriptions	Partial	Module list pages render server-side; module <i>detail</i> pages on App A do not. See P1.
Course Schema potential	N/A	Schema.org/Course is for discoverable public catalogues. These apps are personal and gated; SEO is not a goal.
Titles, durations, pricing	Partly present	Titles & durations: yes. Pricing: not applicable. The duration metadata is well-structured and would be Schema-ready if you ever do publish.

2. Authority & trust

Sub-criterion	Status	Note
Professional, authoritative tone	Strong	Calm, scoped, school-aware language. Avoids hype.
Instructor credentials	N/A	Single-user tool; no instructor surface required.
Partner badges (AWS, CompTIA)	N/A	There's no commercial partnership to display. The A+ alignment is content-level, not certification-level.
Student testimonials	N/A	Personal app; no learner pool. If you ever publish, this would be the place for staff feedback.
"Last updated" date for technical content	Missing	Most important applicable item in this pillar — addressed in P2.

3. Learner UX & clarity

Sub-criterion	Status	Note
Clear learning path	Partial	Modules are well-categorised (Foundations, Networking, Endpoint Support, Identity, Cloud, Operations) but there's no <i>recommended sequence</i> for a learner starting from zero. App A has a <i>/paths</i> route — bring it forward.
Prerequisites visible	Missing	L2 modules (VLANs, MDM/Intune/GPO) sit alongside L1 modules with no prerequisite hint. Add a tiny "Best after: ..." line on L2 cards.

End goal visible	Strong	"Question-first, scenario-based growth across diagnosis, explanation, escalation, documentation, and practical support outputs" — clear and well-scoped.
Scannable syllabus	Mostly	Module list cards scan well. Tag-chip styling regression (P3) is the main blocker.
Obvious CTA / next action	Partial	App B's "What should I study next" section does this well. App A's 11 equally weighted Quick Actions don't. See P2.

4. Content accuracy

I could not assess the actual technical content of modules end-to-end — module bodies on App A are absent in server responses (P1), and a deep technical review of all 36 App B modules is outside the scope of a single audit pass. Spot checks below are based on titles, subtitles, tags, and the diagnostic prompts visible on the working module pages.

Sub-criterion	Status	Note
Outdated terminology	Clean (sample)	No references to Server 2012, SBS, “Office 365”-as-brand, classic Skype, or other dated terms. Entra/Intune/M365/Jamf are current, correct names.
A+ alignment	Reasonable	Modules tagged 220-1201 (current Core 1 objective code — correct as of audit). Worth verifying nothing references the older 220-1101 cycle.
Australian/school context accuracy	Strong	Sentral, OurDCS/Schoolbox, Parent Portal, ViewBoard, PaperCut/Follow-Me, Jamf for iPad — all real and correctly named tools.
Readability without condescension	Mostly strong	Subtitles explain the “why” in plain English. The “without pretending” phrasing (P3) is the only consistent friction point.
Cross-app content drift	Watch	Same module slug (e.g. <i>ports-and-protocols</i>) appears on both apps with different durations (20 vs 20 min — matches in this case, but several others differ slightly). With two apps live, content drift is a real risk.

5. Mobile & accessibility

Sub-criterion	Status	Note
Viewport meta	Present	Both apps set <i>width=device-width</i> , <i>initial-scale=1</i> . App A module pages also set <i>maximum-scale=1</i> and PWA meta tags — caution: maximum-scale=1 prevents pinch-zoom and is an accessibility regression . Drop it.
Theme colour	Set on App A	<i>theme-color</i> = #0f172a — consistent with the brand. Mirror onto App B for parity.
Tables on small screens	Could not assess	Tabular data not visible in fetched markup. Run a manual phone test on Modules and PD Log; long tag chips and progress strings are likely overflow risks.
Code snippets & fixed-width content	None observed	Sampled pages had no code blocks. If you add CLI examples (likely under A+ Core, networking, scripting), use <i>overflow-x: auto</i> with monospace and a min font size.
Top-nav on small screens	Concern	App A’s 16–27 inline nav links will wrap into a wall of text on a phone. Collapse to a hamburger or grouped sections (see P2).

Heading hierarchy & landmarks	Partial	H1 used appropriately. Need to verify nav has <nav> landmarks, main content has <main>, and the duplicated nav (P2) isn't announced twice to AT.
-------------------------------	---------	--

What's genuinely good

A balanced audit names what works, not just what doesn't.

✓ **Operational priority is unmissable.**

"Tickets, walk-ups, calls, and Paul's instructions take priority" appears on every page on both apps. For a tool whose biggest risk is being used at the wrong moment, this is exactly the right design choice. It's also a quiet way of building trust with your supervisor that the tool isn't pulling you away from real work.

✓ **Privacy-first ethos baked into the surface.**

Both apps repeatedly remind the user not to enter real student, staff, parent, credential, or network detail. For a school context this is critical; the tool is structurally less risky than a generic note-taking app for the same purpose.

✓ **Question-first / diagnostic-first pedagogy.**

App B's "Attempt the prompt first, then reveal the guided answer" pattern is well-supported by retrieval-practice research and is rare in DIY learning tools. It's a stronger pedagogical bones than most paid PD platforms.

✓ **DCS-specific naming and scenarios.**

Modules name the actual tools (Sentral, OurDCS, Schoolbox, Parent Portal, ViewBoard, Jamf, PaperCut/Follow-Me) and the actual workflows (access keys, reporting periods, BYOD/COPE). The content feels real because it *is* real.

✓ **Scenarios with multi-step decisions and clearly wrong options.**

The "No internet in classroom" scenario gives three options where two are obviously wrong (open advanced adapter settings; declare an outage). That's the right shape for L1 decision practice — force a clear classification under pressure.

✓ **Evidence Pack & PD Log = audit trail.**

The combination of PD Log (what you did) and Evidence Pack (what you can show) is exactly what self-directed PD needs to count toward formal review or career conversations. This is genuinely missing from most commercial offerings, which prioritise certificates over evidence.

✓ **Honesty about scope.**

Phrases like "Level 1 safe", "before MDM escalation", and "evidence capture" consistently locate the learner inside their actual remit. The tool is teaching *scope* alongside content — a rare and valuable choice.

Recommended next actions

A pragmatic priority order, sized for quiet PD windows.

#	Action	Effort	Impact
1	Pick the canonical app (suggest App B) and decide what becomes of App A — redirect, take offline, or rebrand as “Sandbox”.	S	High
2	Diagnose App A's empty module detail pages in a real browser with DevTools open. Either fix or formally retire those routes.	M	High
3	Add a “Last reviewed” date to every module on the canonical app. Set a 6-month soft expiry on the badge colour.	S	High
4	Replace empty-state metric grid with task-shaped first actions on Day 1; show real metrics only after they have content.	S	Medium
5	Drop <i>maximum-scale=1</i> from the App A viewport meta — small accessibility win.	XS	Medium
6	Group navigation into Learn / Practice / Track / Reference / Admin on whichever app is canonical.	M	Medium
7	Sweep module subtitles for “without pretending / overpromising / ever” phrasing and rewrite as positive scope.	S	Low
8	Fix tag-chip rendering so tags read as discrete words.	XS	Low
9	Add prerequisite hints to L2 modules (“Best after: DNS, DHCP, Gateway, IP Basics”).	S	Low
10	Manual phone-screen QA pass on Modules list, dashboard, and PD Log forms.	S	Low

Effort key

XS = under 30 minutes · **S** = under a focus block (20–90 min) · **M** = a couple of focus blocks · **L** = a weekend project. Sized for quiet windows, not after-hours marathons.

Closing notes for Josh

Both of these apps reflect real care. The pedagogical choices — question-first, scope-aware, evidence-led, privacy-first — are stronger than what most commercial PD platforms do, and the DCS-specific naming makes the content immediately useful rather than generic.

The biggest risk isn't the content; it's the parallel deployment. Two production apps with the same name fragments the data the tools depend on (streaks, due reviews, evidence) and doubles the maintenance surface. Picking one and being deliberate about the other is the single highest-leverage move on this list.

Everything else — the empty-state metrics, the navigation surface, the “last reviewed” dates — is small, focused, doable in a few quiet windows. None of it requires rebuilding what you've already done; most of it is sanding the corners of something that's already pointing the right direction.

Audit prepared 8 May 2026 against the public Vercel deployments of both apps as of that date. Findings reflect what was visible in server responses; some content rendered only after client-side JavaScript may differ from the static markup reviewed here. A second pass with browser DevTools open would confirm or revise the App A module-detail finding (P1).